

```
In [1]: import pandas as pd;
import numpy as np;
import scipy.stats as stats;
import matplotlib.pyplot as plt;
import seaborn as sns;
import sklearn;
from sklearn import linear_model;
from sklearn.metrics import mean_squared_error;
from statsmodels.formula.api import ols;
import statsmodels.api as sm;
from statsmodels.stats.outliers_influence import variance_inflation_factor;
from statsmodels.tools.tools import add_constant;
from platform import python_version
```

```
In [2]: df = pd.read_csv('churn_clean.csv')
```

```
In [3]: df["Churn"] = df['Churn'].map({'No':0, 'Yes':1})
df['Churn'].value_counts()
```

```
Out[3]: 0    7350
1     2650
Name: Churn, dtype: int64
```

```
In [4]: dfmlr = df[['MonthlyCharge', 'Churn', 'Children', 'Income', 'Bandwidth_GB_Year', 'Tenure']]
```

```
In [5]: #C2 summary statistics
dfmlr.describe()
```

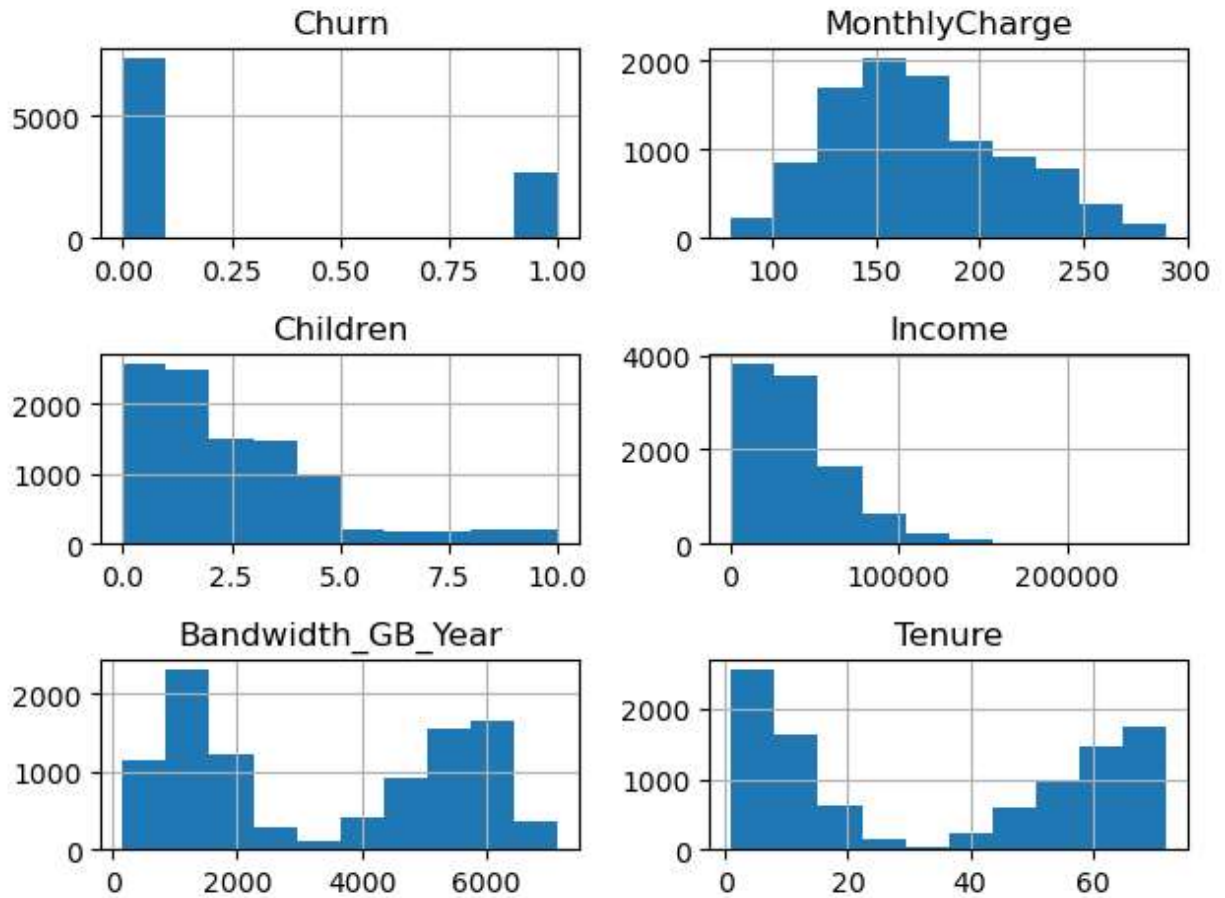
```
Out[5]:
```

	MonthlyCharge	Churn	Children	Income	Bandwidth_GB_Year	Tenure
<b>count</b>	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000	10000.000000
<b>mean</b>	172.624816	0.265000	2.0877	39806.926771	3392.341550	34.526188
<b>std</b>	42.943094	0.441355	2.1472	28199.916702	2185.294852	26.443063
<b>min</b>	79.978860	0.000000	0.0000	348.670000	155.506715	1.000259
<b>25%</b>	139.979239	0.000000	0.0000	19224.717500	1236.470827	7.917694
<b>50%</b>	167.484700	0.000000	1.0000	33170.605000	3279.536903	35.430507
<b>75%</b>	200.734725	1.000000	3.0000	53246.170000	5586.141370	61.479795
<b>max</b>	290.160419	1.000000	10.0000	258900.700000	7158.981530	71.999280

```
In [6]: #get missing values
dfmlr.isna().sum()
```

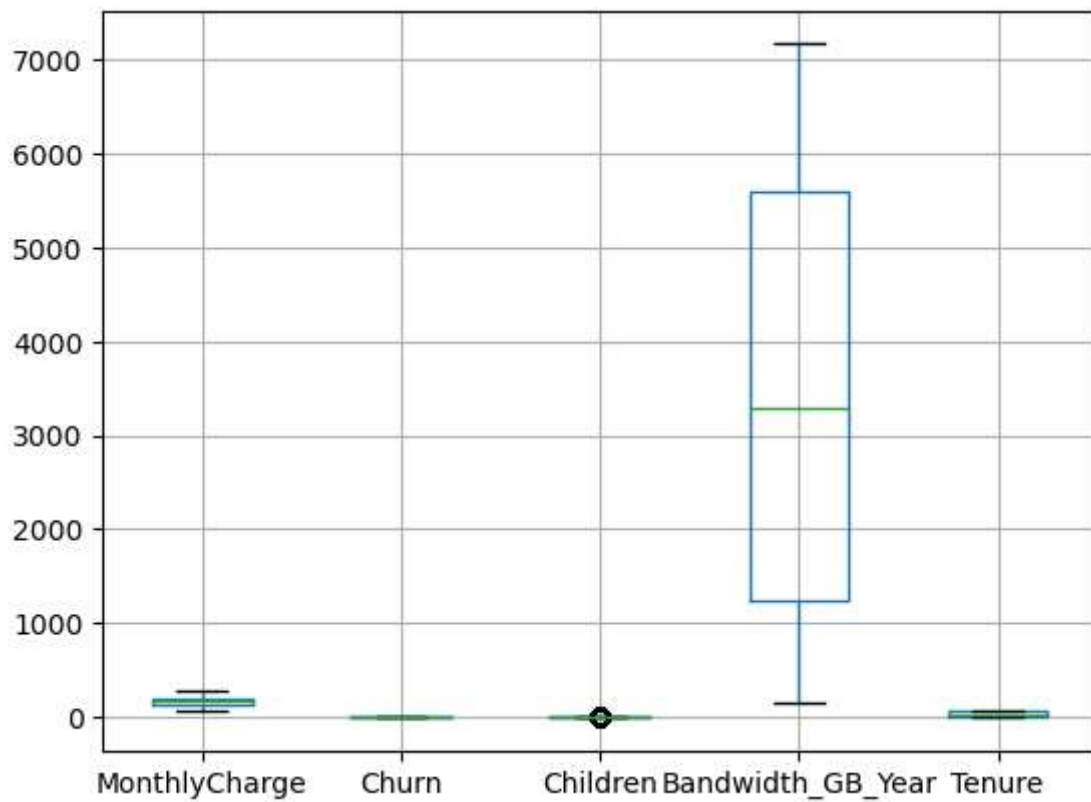
```
Out[6]: MonthlyCharge    0
Churn                  0
Children               0
Income                 0
Bandwidth_GB_Year     0
Tenure                 0
dtype: int64
```

```
In [7]: df[['Churn', 'MonthlyCharge', 'Children', 'Income', 'Bandwidth_GB_Year', 'Tenure']].hist()  
plt.tight_layout()
```




```
In [8]: dfmlr.boxplot(column=['MonthlyCharge', 'Churn', 'Children', 'Bandwidth_GB_Year', 'Tenure'])
```

```
Out[8]: <AxesSubplot:>
```



```
In [9]: dfmlr.boxplot(column=['Income'])
```

```
Out[9]: <AxesSubplot:>
```




```
In [10]: dfmlr.to_excel('MLR_Churn_clean.xlsx')
dfmlr.head()
```

```
Out[10]:
```

	MonthlyCharge	Churn	Children	Income	Bandwidth_GB_Year	Tenure
0	172.455519	0	0	28561.99	904.536110	6.795513
1	242.632554	1	1	21704.77	800.982766	1.156681
2	159.947583	0	4	9609.57	2054.706961	15.754144
3	119.956840	0	1	18925.23	2164.579412	17.087227
4	149.948316	1	0	40074.19	271.493436	1.670972


```
In [11]: sns.scatterplot(x='Income',
                        y='Bandwidth_GB_Year', data=dfmlr)
```

```
Out[11]: <AxesSubplot:xlabel='Income', ylabel='Bandwidth_GB_Year'>
```




```
In [12]: sns.scatterplot(x='MonthlyCharge',
                        y='Bandwidth_GB_Year', data=dfmlr)
```

```
Out[12]: <AxesSubplot:xlabel='MonthlyCharge', ylabel='Bandwidth_GB_Year'>
```




```
In [13]: sns.regplot(x="Children",  
                    y="Bandwidth_GB_Year",  
                    data=dfmlr);
```






```
In [14]: sns.regplot(x="Tenure",  
                    y="Bandwidth_GB_Year",  
                    data=dfmlr);
```



```
In [15]: sns.scatterplot(x='MonthlyCharge',  
                        y='Bandwidth_GB_Year', data=dfmlr)
```

```
Out[15]: <AxesSubplot:xlabel='MonthlyCharge', ylabel='Bandwidth_GB_Year'>
```



In [16]: *#C1 VIF - Checking for multicollinearity*

```
X = dfmlr[['Churn', 'MonthlyCharge', 'Children', 'Income', 'Tenure']]
vif = pd.DataFrame()
vif['feature'] = X.columns
vif["VIF"] = [variance_inflation_factor(X.values, i)
               for i in range(len(X.columns))]
print(vif)
```

	feature	VIF
0	Churn	2.160821
1	MonthlyCharge	7.223742
2	Children	1.857087
3	Income	2.721804
4	Tenure	3.565599

In [17]: *#MLR using statsmodels*

```
X = dfmlr[['Churn', 'MonthlyCharge', 'Children', 'Income', 'Tenure']]
y = dfmlr['Bandwidth_GB_Year']
X = sm.add_constant(X)
mod1 = sm.OLS(y, X).fit()

print(mod1.summary())
```

# OLS Regression Results

Dep. Variable:	Bandwidth_GB_Year	R-squared:	0.989			
Model:	OLS	Adj. R-squared:	0.989			
Method:	Least Squares	F-statistic:	1.720e+05			
Date:	Thu, 05 Jan 2023	Prob (F-statistic):	0.00			
Time:	12:56:43	Log-Likelihood:	-68750.			
No. Observations:	10000	AIC:	1.375e+05			
Df Residuals:	9994	BIC:	1.376e+05			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	-56.0022	10.998	-5.092	0.000	-77.561	-34.444
Churn	127.2987	6.706	18.981	0.000	114.153	140.445
MonthlyCharge	2.7724	0.060	46.007	0.000	2.654	2.890
Children	31.8840	1.091	29.221	0.000	29.745	34.023
Income	9.688e-05	8.31e-05	1.166	0.244	-6.6e-05	0.000
Tenure	82.9982	0.104	799.137	0.000	82.795	83.202
=====						
Omnibus:	2216.683	Durbin-Watson:	1.990			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	639.489			
Skew:	0.383	Prob(JB):	1.37e-139			
Kurtosis:	2.026	Cond. No.	2.29e+05			
=====						

## Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.29e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [18]: #reduced MLR model
X = dfmlr[['MonthlyCharge', 'Children', 'Tenure']]
y = dfmlr['Bandwidth_GB_Year']
X = sm.add_constant(X)
modl = sm.OLS(y, X).fit()

print(modl.summary())
```



# OLS Regression Results

=====						
Dep. Variable:	Bandwidth_GB_Year	R-squared:	0.988			
Model:	OLS	Adj. R-squared:	0.988			
Method:	Least Squares	F-statistic:	2.767e+05			
Date:	Thu, 05 Jan 2023	Prob (F-statistic):	0.00			
Time:	12:56:43	Log-Likelihood:	-68928.			
No. Observations:	10000	AIC:	1.379e+05			
Df Residuals:	9996	BIC:	1.379e+05			
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	-66.5773	10.646	-6.254	0.000	-87.446	-45.708
MonthlyCharge	3.2580	0.056	58.673	0.000	3.149	3.367
Children	31.8156	1.111	28.649	0.000	29.639	33.992
Tenure	81.9695	0.090	909.025	0.000	81.793	82.146
=====						
Omnibus:	3510.726	Durbin-Watson:	1.988			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	723.629			
Skew:	0.389	Prob(JB):	7.35e-158			
Kurtosis:	1.936	Cond. No.	809.			
=====						

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [19]: residuals = mod1.resid
print(residuals)
```

```
0      -147.763598
1       -49.554067
2       181.560974
3       407.896116
4      -287.422847
...
9995    371.094440
9996   -44.144476
9997  -246.438124
9998  -147.488322
9999   -9.042070
Length: 10000, dtype: float64
```


```
In [20]: y_pred=mod1.predict(X)
y_pred
```

```
Out[20]: 0      1052.299708
1       850.536833
2       1873.145987
3       1756.683296
4        558.916283
...
9995    6140.158161
9996    5740.096286
9997    4405.743923
9998    6615.945074
9999    5866.628237
Length: 10000, dtype: float64
```

```
In [21]: mean_squared_error(y,y_pred)
```

```
Out[21]: 56825.44473679924
```

```
In [22]: plt.scatter(modl.predict(), residuals);  
plt.axhline(0, color='red')  
plt.xlabel('Predicted Values');  
plt.ylabel('Residuals');  
plt.xlim([1,1000]);
```



```
In [23]: python_version()
```

```
Out[23]: '3.9.13'
```

```
In [24]: !jupyter --version
```

Selected Jupyter core packages...

IPython	: 7.31.1
ipykernel	: 6.15.2
ipywidgets	: 7.6.5
jupyter_client	: 7.3.4
jupyter_core	: 4.11.1
jupyter_server	: 1.18.1
jupyterlab	: 3.4.4
nbclient	: 0.5.13
nbconvert	: 6.4.4
nbformat	: 5.5.0
notebook	: 6.4.12
qtconsole	: 5.2.2
traitlets	: 5.1.1